

ZML RAG Re-ranking System Data Flow

The ZML RAG Re-ranking System implements a multi-agent architecture for enhanced information retrieval and response generation. At its core, the system employs multiple specialized agents working in concert: Query Processing Agents analyze and expand queries using domain-specific knowledge, Retrieval Agents perform parallel semantic and keyword searches, and Ranking Agents evaluate results based on relevance, diversity, and authority. The re-ranking process leverages a multi-stage pipeline where responses are scored and reordered using weighted criteria and contextual relevance. This is integrated with a vector store for efficient knowledge retrieval and a memory system for maintaining conversation context. A key innovation is the implementation of parallel processing paths where multiple agents can simultaneously evaluate different aspects of the response, with their outputs being combined through a rank fusion algorithm that considers both individual scores and inter-document relationships. The system's effectiveness is enhanced by its use of BioBERT for domain-specific embeddings and GPT-4 for natural language understanding, ultimately delivering contextually appropriate, diverse, and authoritative responses with proper citations.

https://github.com/tmlrnc/LLM/blob/main/ZML/zml_rag_4.py

1. Data Flow Components:

- User Interface Layer: Handles query input and response output
- Core System: Main processing components
- Vector Store Layer: Knowledge base and embeddings
- Agent Layer: Specialized processing agents
- Memory System: Conversation and context storage
- LLM Layer: Language model processing

2. Key Processes:

- Query Processing
- Agent Selection
- Context Management
- Response Ranking
- Memory Management

3. Data Flows:

- User query → Query Processor
- Query → Vector Store
- Context → Agents
- Agent outputs → Response Ranker
- Final response → User

4. Control Flows:

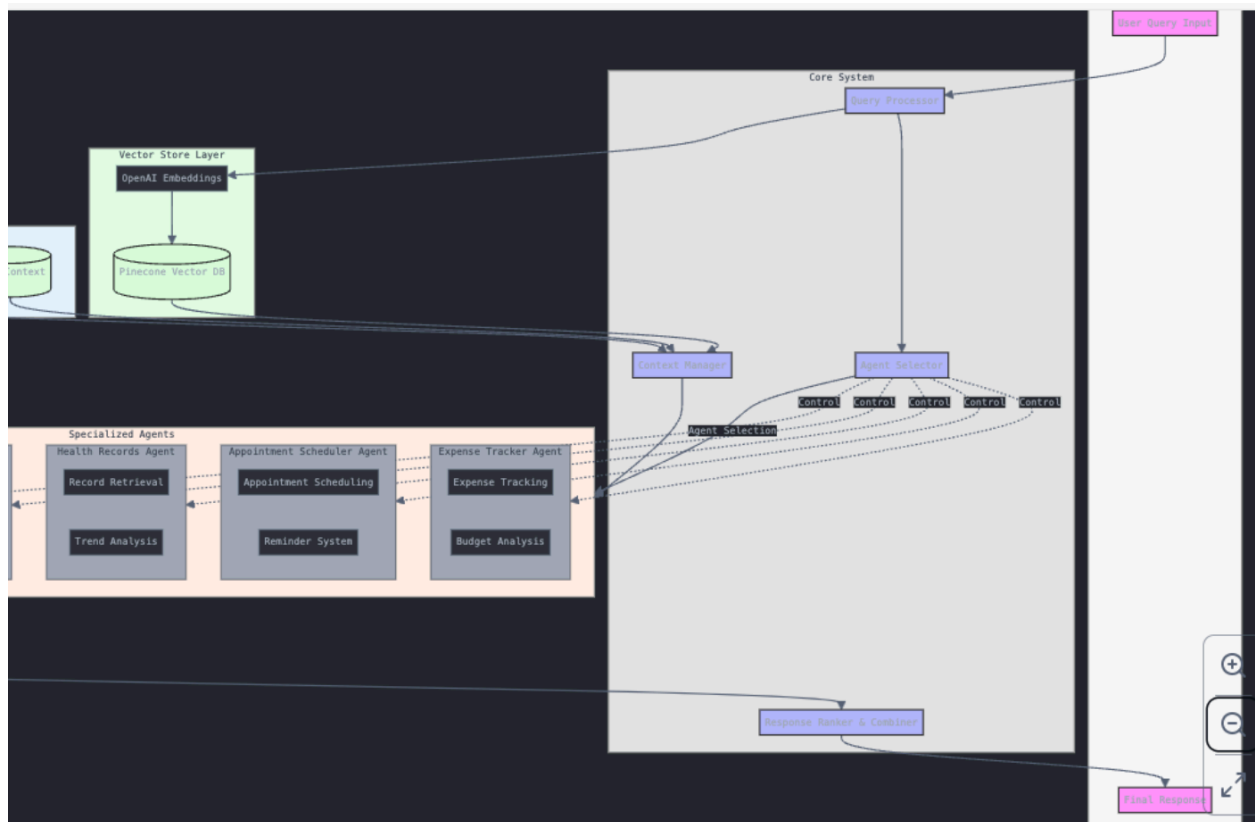
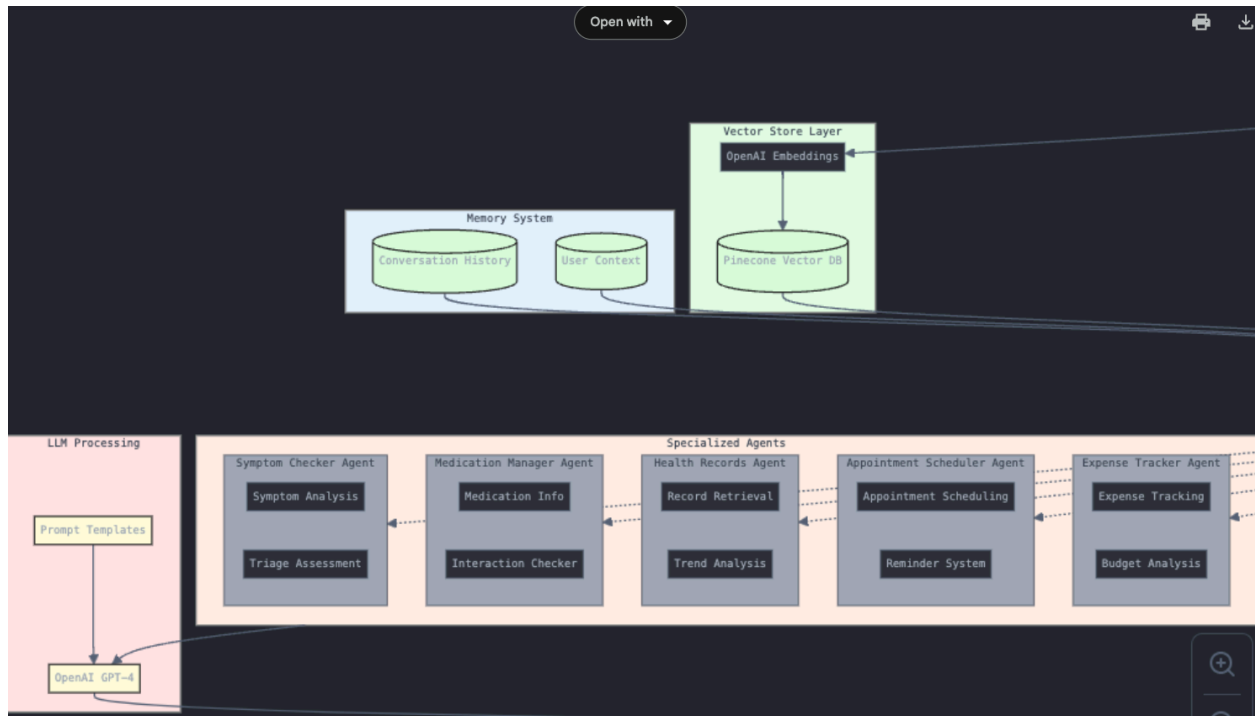
- Agent Selector → Specialized Agents
- Context Manager → Agents
- Memory System → Context Manager

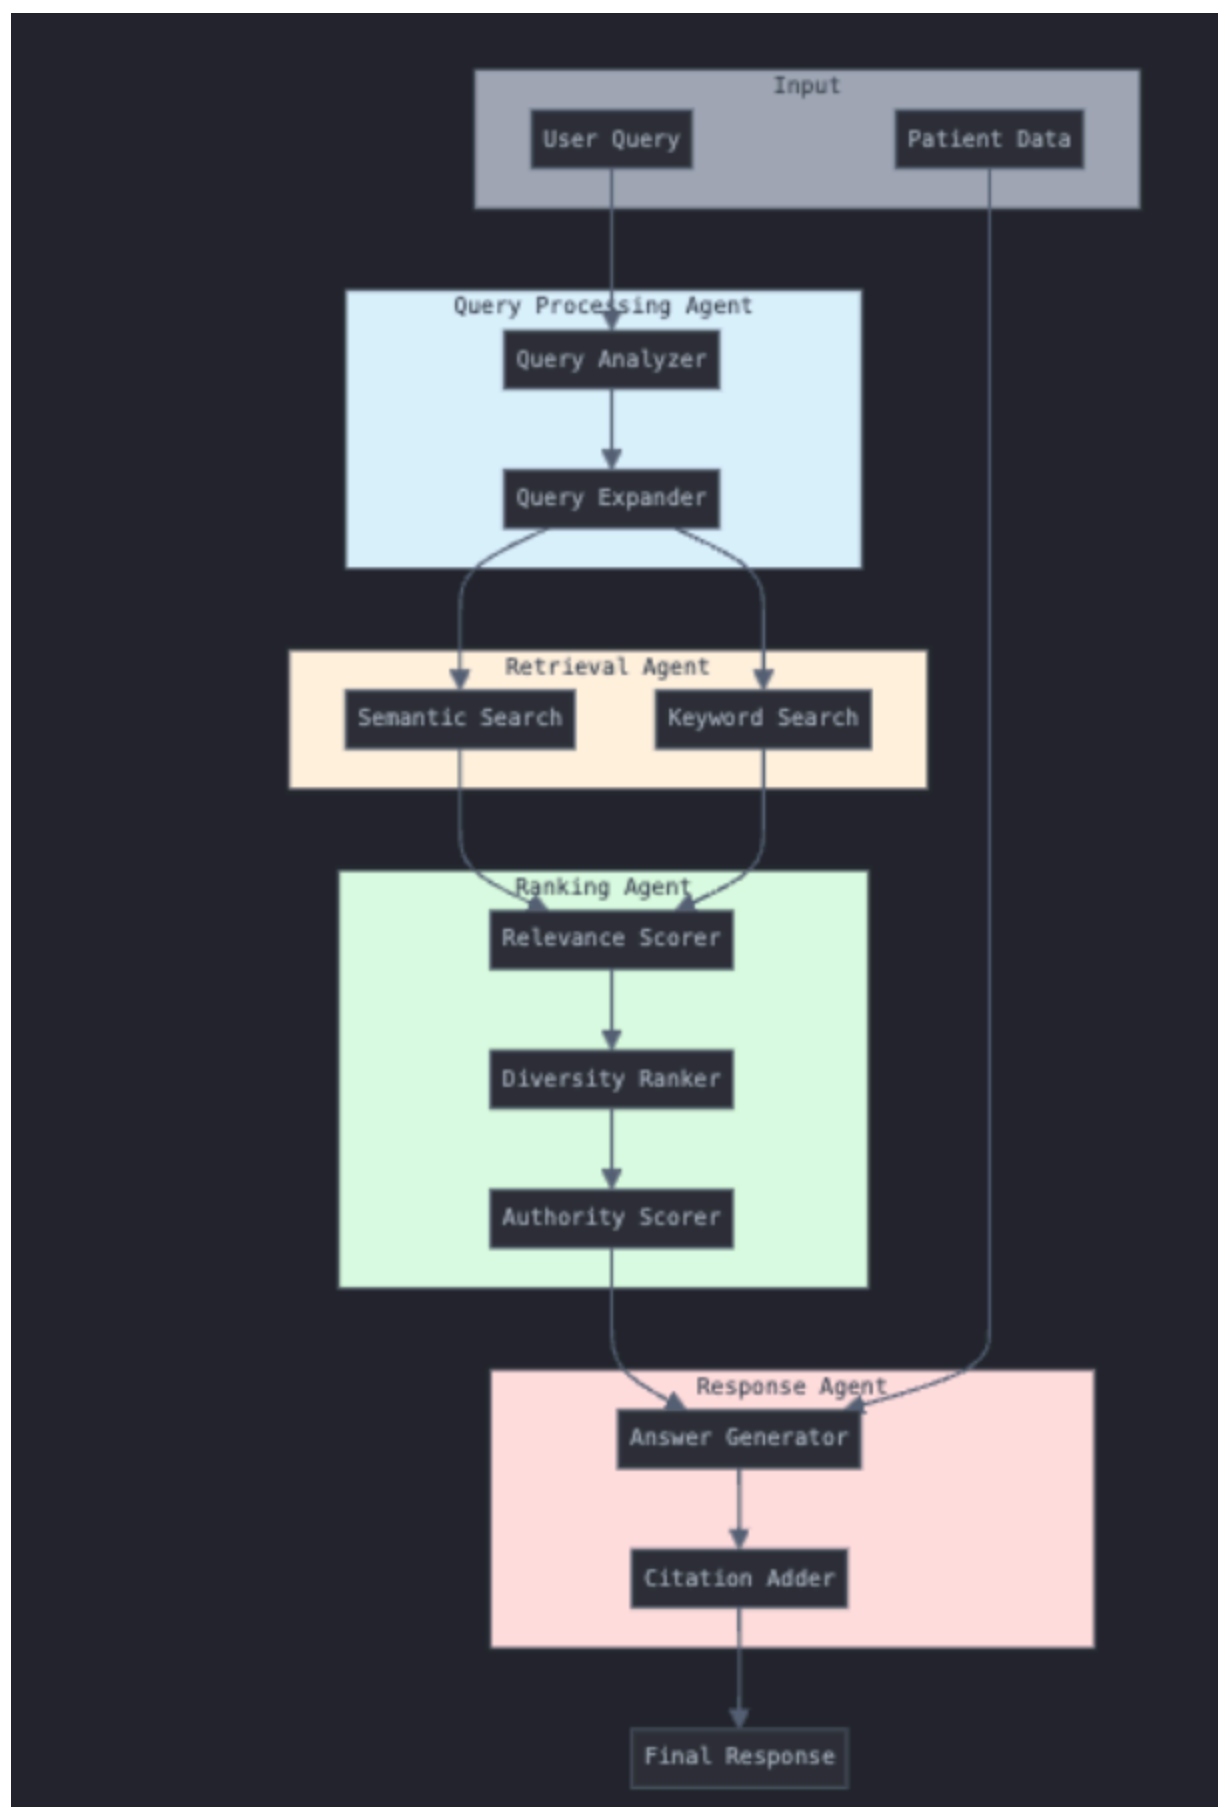
The system follows these steps:

1. User submits query
2. Query Processor analyzes and embeds query
3. Agent Selector chooses relevant agents
4. Agents process query with context
5. Responses are ranked and combined
6. Final response is returned to user
7. Input Stage:
 - a. User Query: Initial question from the user
 - b. Patient Data: Contextual medical information about the patient
8. Query Processing Agent:
 - a. Query Analyzer: Examines query intent and context
 - b. Query Expander: Enriches query with relevant medical terms
9. Retrieval Agent:
 - a. Semantic Search: Performs embedding-based similarity search
 - b. Keyword Search: Traditional term-based search Both methods run in parallel to maximize coverage
10. Ranking Agent:
 - a. Relevance Scorer: Evaluates document relevance to query
 - b. Diversity Ranker: Ensures varied information sources
 - c. Authority Scorer: Weights results based on source credibility
11. Response Agent:
 - a. Answer Generator: Synthesizes information with patient context
 - b. Citation Adder: Includes reference sources in response

The system uses specialized agents at each stage, with BioBERT for medical text encoding and GPT-4 for natural language understanding and generation. Each stage feeds into the next, with the final output being a contextualized medical response with citations.

Architecture Diagram





1. Knowledge Graph Schema :
 - Shows the relationships between different medical entities
 - Includes patients, conditions, medications, symptoms, etc.
 - Defines key relationships like HAS_CONDITION, TAKES, INTERACTS_WITH
2. Neo4j Integration Code:
 - **MediMateKnowledgeGraph** class for Neo4j operations:
 - Schema setup with constraints
 - CRUD operations for medical data
 - Specialized queries for medical knowledge
 - Integration with existing agent infrastructure
3. Key Features:
 - Medical history tracking
 - Medication interaction checking
 - Symptom-condition relationships
 - Patient record management
 - Treatment and insurance tracking

